

NCCL Testing Over RoCE Fabric - Complete Tutorial

Overview

This guide covers setting up and running NCCL (NVIDIA Collective Communications Library) tests between bare metal GPU nodes using RoCE (RDMA over Converged Ethernet) fabric for high-performance inter-node communication.

Prerequisites

- Multiple bare metal servers with NVIDIA GPUs

Installation

None

```
apt install -y git
rm -rf nccl-tests
git clone https://github.com/NVIDIA/nccl-tests.git
cd nccl-tests
make -j16 MPI=1 MPI_HOME=/usr/lib/x86_64-linux-gnu/openmpi/
```

1. Hardware and Network Setup

Network Infrastructure Requirements

- **Switch Configuration:** Enable PFC (Priority Flow Control) and ECN (Explicit Congestion Notification)
- **Cabling:** Use appropriate cables (typically 25GbE, 50GbE, or 100GbE)
- **Network Topology:** Ensure low-latency, high-bandwidth connectivity between nodes

Verify GPU and Network Hardware

Shell

Check GPUs

nvidia-smi

Check network interfaces

ip link show

ibv_devices *# Should show RDMA devices*

Check RoCE capability

show_gids | grep -i roce

Method 1: Use rdma command (more modern approach)

rdma link show

Method 2: Use ibv_devinfo to check port capabilities

ibv_devinfo -v | grep -i roce

2. Software Installation

Install Required Packages

Shell

Update system

sudo apt update && sudo apt upgrade -y

Install essential packages (optional: already in the system)

sudo apt install -y build-essential cmake git

sudo apt install -y libibverbs-dev librdmacm-dev

sudo apt install -y rdma-core perftest

Install NVIDIA drivers (if not already installed)

wget

https://developer.download.nvidia.com/compute/cuda/repos/ubuntu2004/x86_64/cuda-keyring_1.1-1_all.deb

sudo dpkg -i cuda-keyring_1.1-1_all.deb

```
sudo apt update
sudo apt install -y cuda-drivers cuda-toolkit-12-4
```

Install NCCL

Shell

Download and install NCCL

```
wget
https://developer.download.nvidia.com/compute/redist/nccl/v2.19.3
/nccl_2.19.3-1+cuda12.0_x86_64.txz
tar -xvf nccl_2.19.3-1+cuda12.0_x86_64.txz
sudo cp -r nccl_2.19.3-1+cuda12.0_x86_64/* /usr/local/
sudo ldconfig
```

Verify installation

```
pkg-config --modversion nccl
```

Install NCCL Tests (if not installed in above step)

Shell

Clone NCCL tests repository

```
git clone https://github.com/NVIDIA/nccl-tests.git
cd nccl-tests
```

Build NCCL tests

```
make MPI=1 MPI_HOME=/usr/lib/x86_64-linux-gnu/openmpi
CUDA_HOME=/usr/local/cuda NCCL_HOME=/usr/local
```

Quick Setup

1. Install NCCL Tests

Shell

Install prerequisites

```
apt install -y git openmpi-bin openmpi-common libopenmpi-dev
```

Build NCCL tests

```
cd ~
```

```
rm -rf nccl-tests
```

```
git clone https://github.com/NVIDIA/nccl-tests.git
```

```
cd nccl-tests
```

```
make -j16 MPI=1 MPI_HOME=/usr/lib/x86_64-linux-gnu/openmpi/
```

2. Find Your RoCE Interfaces

Shell

Find interfaces with RDMA devices (these are your RoCE interfaces)

```
for iface in $(ls /sys/class/net/); do
```

```
    rdma_dev=$(ls /sys/class/net/$iface/device/infiniband/
```

```
2>/dev/null)
```

```
    if [ -n "$rdma_dev" ]; then
```

```
        echo "RoCE Interface: $iface → RDMA device: $rdma_dev"
```

```
    fi
```

```
done
```

Alternative: Use ibdev2netdev to see mappings

```
ibdev2netdev
```

Sample Output:

None

RoCE Interface: enp156s0np0 → RDMA device: mlx5_6

3. Check GPU-NIC Topology

Shell

See which NICs are closest to which GPUs

```
nvidia-smi topo -m
```

Single Node Test

Create Test Script

Shell

```
cat > ~/test_single_node.sh << 'EOF'
```

```
#!/bin/bash
```

```
echo "=== Single Node NCCL Test ==="
```

```
# Get GPU count
```

```
GPU_COUNT=$(nvidia-smi -L | wc -l)
```

```
echo "Testing with $GPU_COUNT GPUs"
```

```
cd ~/nccl-tests
```

```
# Run the test (target: >400 GB/s bus bandwidth)
```

```
./build/all_reduce_perf -b 1G -f 2 -e 1G -g $GPU_COUNT
```

```
EOF
```

```
chmod +x ~/test_single_node.sh
```

Run Single Node Test

Shell ▾

```
~/test_single_node.sh
```

Expected Output: Look for **busbw** > 400 GB/s in the results.

Sample Output:

Shell

```
size      count      type  redop  root    time  algbw  busbw #wrong
time  algbw  busbw #wrong
#      (B)    (elements)
(us) (GB/s) (GB/s)
1073741824 268435456 float sum -1 5110.3 210.11 367.70 0
5112.6 210.02 367.53 0
# Out of bounds values : 0 OK
# Avg bus bandwidth : 367.616
#
# Collective test concluded: all_reduce_perf
```

Multi-Node Test

Create Multi-Node Test Script

Shell

```
cat > ~/test_multi_node.sh << 'EOF'
#!/bin/bash

# Usage: ./test_multi_node.sh "ip1,ip2"
HOSTS=$1

if [ -z "$HOSTS" ]; then
    echo "Usage: $0 \"ip1,ip2,ip3\""
    echo "Example: $0 \"192.168.1.10,192.168.1.11\""
    exit 1
```

```

fi

# Parse hosts
IFS=',' read -r -a ip_array <<< "$HOSTS"
NP=$(echo "${ip_array[@]}" | tr ' ' '\n' | sort -u | wc -l)

# Assuming 8 GPUs per node (adjust as needed)
NEW_NP=$(( NP * 8 ))
NEW_HOSTS=$(echo $HOSTS | sed 's/\([0-9\.]*\)/\1:8/g')

echo "Testing $NP nodes, $NEW_NP total GPUs"

cd ~/nccl-tests

# Run multi-node test
mpirun --bind-to none --mca pml ucx --allow-run-as-root \
    -x NCCL_CROSS_NIC=0 \
    -np $NEW_NP -H ${NEW_HOSTS} \
    ./build/all_reduce_perf -b 1G -f 2 -e 1G
EOF

chmod +x ~/test_multi_node.sh

```

Run Multi-Node Test

```

Shell

# Replace with your actual node IPs
~/test_multi_node.sh "10.45.169.116,10.45.169.109"

```

Expected Output: Look for **busbw** > 400 GB/s across nodes.

Troubleshooting

Abrupt test termination

An abrupt test termination is usually associated with hardware problems, like a misconfiguration or a bad cable, so it is worth checking the following:

- all GPU NICs are up
- they are all on the same VLAN
- they are plugged in the right leaf switches

These are example outputs associated to good nodes:

Interfaces are up =>

None

Currently back-end network is addressed via IPv6, this may be subject to change to enable PFC.

ip -br a

lo	UNKNOWN	127.0.0.1/8 ::1/128
eno8303	DOWN	
eno8403	DOWN	
enp26s0np0	UP	
fe80::a288:c2ff:fee4:446a/64		
enp27s0f0np0	UP	
enp27s0f1np1	UP	
enp60s0np0	UP	
fe80::a288:c2ff:fee3:638a/64		
enp77s0np0	UP	
fe80::a288:c2ff:feda:31ea/64		
enp94s0np0	UP	
fe80::a288:c2ff:feda:d712/64		
enp156s0np0	UP	
fe80::a288:c2ff:feda:d6c2/64		
enp157s0f0np0	UP	
enp157s0f1np1	UP	


```
enp188s0np0      UP
fe80::5aa2:e1ff:fe11:5ffc/64
enp204s0np0      UP
fe80::a288:c2ff:fee4:4322/64
enp220s0np0      UP
fe80::a288:c2ff:feda:d6e2/64
bond0             UP           10.45.169.109/24
fe80::9e63:c0ff:fef6:1eb2/64
```

Interfaces are on the same VLAN =>

None

```
# lldpctl | grep -i VLAN
VLAN:      108, pvid: yes vlan-108
VLAN:      108, pvid: yes vlan-108
VLAN:      108, pvid: yes vlan-108
VLAN:      108, pvid: yes vlan-108
VLAN:      108, pvid: yes vlan-108
VLAN:      108, pvid: yes vlan-108
VLAN:      108, pvid: yes vlan-108
VLAN:      108, pvid: yes vlan-108
VLAN:      108, pvid: yes vlan-108
VLAN:      108, pvid: yes vlan-108
VLAN:      108, pvid: yes vlan-108
VLAN:      108, pvid: yes vlan-108
```

Interfaces are ordered correctly on leaf switches (compare output on both nodes - order must match) =>

None

```
# lshw -c network -businfo | grep X-7 | awk '{print $2}' |
xargs lldpctl | grep SysName
SysName:      am4-t0-g1-leaf07
```

```
SysName:      am4-t0-g1-leaf04
SysName:      am4-t0-g1-leaf03
SysName:      am4-t0-g1-leaf08
SysName:      am4-t0-g1-leaf06
SysName:      am4-t0-g1-leaf05
SysName:      am4-t0-g1-leaf01
SysName:      am4-t0-g1-leaf02
```

If Test Hangs

Add debug info:

Shell

Add to your multi-node script

```
mpirun --bind-to none --mca pml ucx --allow-run-as-root \
-x NCCL_CROSS_NIC=0 \
-x NCCL_DEBUG=INFO \
-np $NEW_NP -H ${NEW_HOSTS} \
./build/all_reduce_perf -b 1G -f 2 -e 1G
```

Look for messages like:

- **Good:** via **NET/IB/1/GDRDMA** (RDMA working)
- **Bad:** via **NET/IB/1** (falling back to system memory)

Expected Output:

None

```
atl1g3r4u9gpu:146600:146782 [5] transport/net_ib.cc:2244 NCCL WARN NET/IB: Got
completion from peer 10.46.70.21<50854> with status=5 opcode=31245 len=0 vendor
```

```
err 249 (Recv) localGid fe80::9e63:c0ff:feab:b8da
remoteGidsfe80::9e63:c0ff:feab:b9ca hca mlx5_9
atl1g3r4u9gpu:146600:146782 [5] NCCL INFO transport/net.cc:1362 -> 6
atl1g3r4u9gpu:146600:146782 [5] NCCL INFO proxy.cc:728 -> 6
atl1g3r4u9gpu:146600:146782 [5] NCCL INFO proxy.cc:912 -> 6 [Progress Thread]

atl1g3r1u16gpu:496000:496186 [6] transport/net_ib.cc:2244 NCCL WARN NET/IB: Got
completion from peer 10.46.70.22<39226> with status=12 opcode=32224 len=0
vendor err 129 (Recv) localGid fe80::9e63:c0ff:feab:938a
remoteGidsfe80::9e63:c0ff:feab:b6b2 hca mlx5_10
atl1g3r1u16gpu:496000:496186 [6] NCCL INFO transport/net.cc:1362 -> 6
atl1g3r1u16gpu:496000:496186 [6] NCCL INFO proxy.cc:728 -> 6
atl1g3r1u16gpu:496000:496186 [6] NCCL INFO proxy.cc:912 -> 6 [Progress Thread]
```

They indicate that pairs of nodes(Node-A and Node-B), in this case 10.46.70.21 and 10.46.70.22, are having communication issue over the GPU fabric. This can be verified with using ping6 from Node-A:

None

```
ping6 -c 2 -I enp156s0np0 fe80::9e63:c0ff:feab:b9ca
```

Where enp156s0np0 is one of the 8 400G Nic, and fe80::9e63:c0ff:feab:b9ca the IPv6 address of the same NIC on Node-B.

Slow bus bandwidth

In this case, you need to isolate the culprit. The most obvious thing to rule out is software misconfiguration. What happens most often than not is that the GPU ends up talking to the NIC off system memory, so RDMA is not place. In other words you would need to see RDMA in-place and nccl tells you so in this form when you pass -x

NCCL_DEBUG=info to its arguments:

None

```
[...]
```

```
am4g1r34bm1:23858:23875 [0] NCCL INFO Channel 00/0 :  
0[6] -> 1[6] [send] via NET/IB/1/GDRDMA  
  
am4g1r34bm1:23858:23875 [0] NCCL INFO Channel 01/0 :  
0[6] -> 1[6] [send] via NET/IB/1/GDRDMA  
  
am4g1r34bm1:23858:23875 [0] NCCL INFO Channel 02/0 :  
0[6] -> 1[6] [send] via NET/IB/1/GDRDMA  
  
am4g1r34bm1:23858:23875 [0] NCCL INFO Channel 03/0 :  
0[6] -> 1[6] [send] via NET/IB/1/GDRDMA
```

When it does not, the output is more like this one:

None

```
[...]  
  
am4g1r35bm1:1226446:1226462 [0] NCCL INFO Channel 00/0  
: 0[6] -> 1[6] [receive] via NET/IB/1  
  
am4g1r35bm1:1226446:1226462 [0] NCCL INFO Channel 01/0  
: 0[6] -> 1[6] [receive] via NET/IB/1  
  
am4g1r35bm1:1226446:1226462 [0] NCCL INFO Channel 00/0  
: 1[6] -> 0[6] [send] via NET/IB/1  
  
am4g1r35bm1:1226446:1226462 [0] NCCL INFO Channel 01/0  
: 1[6] -> 0[6] [send] via NET/IB/1
```

And this is bad. In this case check that all the nvidia drivers are loaded correctly:

None

```
# lsmod | grep nvid  
  
nvidia_uvm          4956160  0
```

```
nvidia_peermem          16384  0

ib_uverbs                200704  3
nvidia_peermem,rdma_ucm,mlx5_ib

nvidia_drm               122880  0

nvidia_modeset           1355776  1 nvidia_drm

nvidia                   54292480  63
nvidia_uvm,nvidia_peermem,nvidia_modeset

video                    73728  2
dell_wmi,nvidia_modeset
```

`nvidia_peermem` is the module in question. If that is not loaded, than you probably need to rebuild it and install it again. Check what you have:

None

```
# dkms status | grep nv

nvidia/550.90.07, 6.8.0-52-generic, x86_64: installed
```

Then do:

None

```
# Remove nvidia driver

dkms remove nvidia/550.90.07


# Install nvidia driver

dkms install nvidia/550.90.07
```

Once that is done, `modprobe module nvidia_peermem` in and verify that is loaded correctly. You can attempt now to perform a new test, if it is still not in the clear, then performance slowness could be the result of the following:

- bad optics
- bad fiber
- bad NIC (if a card is replaced, remember to reset ETH_MODE and ATS attributes, check the [appendix](#) for more details)
- bad GPU

Check NVIDIA Drivers

```
Shell
# Verify nvidia_peermem is loaded (required for RDMA)
lsmod | grep nvidia_peermem

# If missing, reload drivers
sudo dkms remove nvidia/${nvidia-smi --query-gpu=driver_version
--format=csv,noheader,nounits | head -1}
sudo dkms install nvidia/${nvidia-smi --query-gpu=driver_version
--format=csv,noheader,nounits | head -1}
sudo modprobe nvidia_peermem
```

Test Individual GPU-NIC Pairs

If you get low bandwidth, test each GPU-NIC combination:

```
Shell
cat > ~/test_gpu_nic_pair.sh << 'EOF'
#!/bin/bash

# Usage: ./test_gpu_nic_pair.sh "ip1,ip2" gpu_id nic_id
HOSTS=$1
GPU_ID=$2
NIC_ID=$3
```

```

NEW_HOSTS=$(echo $HOSTS | sed 's/\([0-9\.]*\)/\1:1/g')

echo "Testing GPU$GPU_ID with NIC$NIC_ID"

cd ~/nccl-tests

mpirun --bind-to none --mca pml ucx --allow-run-as-root \
    -x NCCL_CROSS_NIC=0 \
    -x NCCL_DEBUG=info \
    -x CUDA_VISIBLE_DEVICES=$GPU_ID \
    -x NCCL_IB_HCA=mlx5_$NIC_ID \
    -np 2 -H ${NEW_HOSTS} \
    ./build/all_reduce_perf -b 1G -f 2 -e 1G
EOF

chmod +x ~/test_gpu_nic_pair.sh

# Test specific GPU-NIC pairs (based on nvidia-smi topo -m)
~/test_gpu_nic_pair.sh "ip1,ip2" 0 0
~/test_gpu_nic_pair.sh "ip1,ip2" 1 3
# etc.

```

Expected single GPU-NIC bandwidth: ~48 GB/s

Network Verification

Check Interface Status

```

Shell
# Check interfaces are up
ip -br a

# Check VLAN (all interfaces should be on same VLAN)
lldpctl | grep -i VLAN

```

```
# Check switch connectivity
```

```
lshw -c network -businfo | grep ConnectX | awk '{print $2}' |  
xargs lldpctl | grep SysName
```

Test Network Connectivity

Shell

```
# Ping between nodes using RoCE interfaces
```

```
ping6 -c 2 -I enp156s0np0 fe80::target_ipv6_address
```


